

Reinforcement Learning

1. Introduction to reinforcement learning

1.1 Why reinforcement learning:

There are problems where we can formulate them as MDP, but we are unsure of its transition model T and reward function R . If we know its T , and R , then we can either run policy iteration or value iteration to obtain an optimal policy. However, when T/R , or T and R , are unknown, we can utilize reinforcement learning.

1.2 Goal of reinforcement learning:

Agent learns to act optimally, i.e. learn to maximize expected rewards, by receiving feedbacks in the form of rewards.

Remark (Utility). *Agent's utility is defined by the reward function.*

1.3 Types of reinforcement learning:

- **Passive learning:** also referred as 'prediction'. We evaluate a fixed policy, instead of learning good actions.
- **Active learning:** learning to act optimally

Each of the above section, can be further divided into:

- **Model-based learning:** learn the model of MDP, i.e. learn the transition probability model $T(s, a, s')$, and observe $R(s, a, s')$ from experience of (s, a, s') . Compute optimal policy by either value iteration or policy iteration.
- **Model-free learning:** Bypass the need to learn T , R , compute an optimal policy from experience of agent its own, or other policies with a Q-learning table, which records state-action values. Passive model-free learning are TD, MC. Active model-free learning can be classified into:
 - **On-policy learning:** Learn about policy π from experience sampled from π , e.g. SARSA

- **Off-policy learning:** Learn about policy π from experience sampled from μ , e.g. Q-learning

Remark (Optimal Policy). *Optimal policy is deterministic.*

Remark (π, μ) .

1.4 Definitions

- **Model:** We refer model as transition probability model, T .

Example to Illustrate Model-Based vs. Model-Free: Expected Age

Goal: Compute expected age of cs188 students

Known $P(A)$

$$E[A] = \sum_a P(a) \cdot a = 0.35 \times 20 + \dots$$

Without $P(A)$, instead collect samples $[a_1, a_2, \dots, a_N]$

Unknown $P(A)$: “Model Based”

$$\hat{P}(a) = \frac{\text{num}(a)}{N}$$

$$E[A] \approx \sum_a \hat{P}(a) \cdot a$$

Unknown $P(A)$: “Model Free”

$$E[A] \approx \frac{1}{N} \sum_i a_i$$

2. Passive reinforcement learning

In passive learning, the policy is fixed. Therefore, we only learn the utility value of each state if the policy is run from that state. The **goal** is to learn the value function $U^\pi(s)$ from observations when **transition model** $P(s'|s, a)$ and **reward function** $R(s)$ are unknown.

General idea:

- The agent executes a set of trials using the fixed policy π .
- The utility or value for π is $U^\pi(s) = E[\sum_{t=0}^{\infty} \gamma^t R(s_t)]$

Question: how to compute $U^\pi(s)$?

2.1 Model-based leaning:

- Adaptive Dynamic Programming: learns the T, R model, then solves it.
 - Learn the transition probabilities $P(s'|s, \pi(s))$
 - Learn reward function $R(s)$
 - Calculate the value by solving the **Bellman equation** with linear algebra
 - * Can use **modified policy iteration** method of doing k iteration of value updates after each change to model
 - Can learn the model using supervised learning. The pseudo-code below shows learning using maximum likelihood with table representing $P(s'|s, a)$

2.2 Model-free leaning:

- Direct Utility Estimation / Monte Carlo Learning
 - **Expected reward-to-go/return:** In a state, the utility or value is the expected total reward from that state onwards
 - Each trial is treated as providing a sample of this quantity for each state visited
 - Keep a **running average** for each state in a table. In infinitely many trial, sample average will converge to expected value.
 - * Running average:

$$\begin{aligned}
U_k(s) &= \frac{1}{k} \sum_{i=0}^k G_i(s) \\
&= \frac{1}{k} (G_k(s) + (k-1)G_{k-1}U_{k-1}(s)) \\
&= \underbrace{\frac{1}{k} G_k(s)}_{\text{Estimate after k-1 returns}} + \underbrace{\frac{1}{k} (G_k(s) - U_{k-1}(s))}_{\text{Prediction Error}}
\end{aligned} \tag{1}$$

- Monte Carlo Learning is just an instance of supervised learning: each example has state as input and observed return as output
- Advantage: Simple; Each labeled target/return is an unbiased estimate
- Disadvantage:
 - * Need to wait till the end of episode before learning can be done
 - * Variance can be high as a return is a sum of many rewards over the sequence. Exploiting the constraint imposed by the Bellman equation may help reduce variance (\implies Temporal-Difference Learning)
- **Temporal difference** learning exploits more of the Bellman equation constraints than Monte Carlo learning. For a transition from state s to s' , TD learning does: $U^\pi(s) = U^\pi(s) + \alpha \underbrace{(R(s) + \gamma U^\pi(s') - U^\pi(s))}_{\text{TD Error}}$
 - α is the learning rate
 - $R(s) + \gamma U^\pi(s')$ is called the TD target
 - $R(s) + \gamma U^\pi(s') - U^\pi(s)$ is called the TD error
 - Converges to the expected value if α decreases with the number of times the states has been visited, e.g. $\alpha(n) = O(1/n)$

Remark (Model-free). *TD does not need a transition model to do the updates, only the observed transition.*

2.3 n-step TD

- Let $G_{t:t+n} = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots + \gamma^n \bar{V}(S_{t+n})$, where $\bar{V}(S_{t+n})$ is the estimated value at state S_{t+n}
- n-step TD sets $G_{t:t+n}$ as the target for update

2.4 TD(λ)

- $G_t^\gamma = (1-\lambda) \sum_{n=1}^{\infty} \gamma^{n-1} G_{t:t+n}$, an average TD value over different values of n.
- G_t^γ is a weighted sum of n-step returns, where n-th step is weighted by λ^{n-1}
- When $\lambda = 1$, we get Monte Carlo. When $\lambda = 0$, we get TD.

$$\begin{aligned}
G_t^\lambda &= (1-\lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_{t:t+n} \\
&= (1-\lambda) \lambda^0 (R_t + V(s_{t+1})) \\
&\quad + (1-\lambda) \lambda^1 (R_t + R_{t+1} + V(s_{t+2})) \\
&\quad + (1-\lambda) \lambda^2 (R_t + R_{t+1} + R_{t+2} + V(s_{t+3})) \\
&\quad + (1-\lambda) \lambda^3 (R_t + R_{t+1} + R_{t+2} + R_{t+3} + V(s_{t+4})) \\
&\quad + \dots \\
&= (1-\lambda) [R_t(\lambda^0 + \lambda^1 + \lambda^2 + \dots + \lambda^\infty) \\
&\quad + R_{t+1}(\lambda^1 + \lambda^2 + \lambda^3 + \dots + \lambda^\infty) \\
&\quad + R_{t+2}(\lambda^2 + \lambda^3 + \lambda^4 + \dots + \lambda^\infty) \\
&\quad + R_{t+3}(\lambda^3 + \lambda^4 + \lambda^5 + \dots + \lambda^\infty) \\
&\quad + \dots] \\
&\quad + (1-\lambda) [\lambda^0 V(s_{t+1}) + \lambda^1 V(s_{t+2}) + \lambda^2 V(s_{t+3}) + \dots] \\
&= (1-\lambda)(1 + \lambda^1 + \dots + \lambda^\infty) [R_t \lambda^0 + R_{t+1} \lambda^1 + R_{t+2} \lambda^2 + \dots] \\
&\quad + (1-\lambda) [\lambda^0 V(s_{t+1}) + \lambda^1 V(s_{t+2}) + \lambda^2 V(s_{t+3}) + \dots] \\
&= (1-\lambda) \cdot \frac{1}{1-\lambda} \cdot [R_t \lambda^0 + R_{t+1} \lambda^1 + R_{t+2} \lambda^2 + \dots] \\
&\quad + (1-\lambda) [\lambda^0 V(s_{t+1}) + \lambda^1 V(s_{t+2}) + \lambda^2 V(s_{t+3}) + \dots] \\
&= R_t \lambda^0 + R_{t+1} \lambda^1 + R_{t+2} \lambda^2 + R_{t+3} \lambda^3 + \dots \\
&\quad + (1-\lambda) [\lambda^0 V(s_{t+1}) + \lambda^1 V(s_{t+2}) + \lambda^2 V(s_{t+3}) + \dots]
\end{aligned}$$

$$\begin{aligned}
\text{if } \lambda \rightarrow 0 \quad G_t^\lambda &= R_t + V(s_{t+1}) && (1\text{-step TD}) \\
\text{if } \lambda \rightarrow 1 \quad G_t^\lambda &= R_t + R_{t+1} + R_{t+2} + \dots && (\text{Monte Carlo})
\end{aligned}$$

2.5 Summary

TD vs. ADP

- ADP learns the model then solves for $U^\pi(s)$ for each state
- TD/MC does not need a model. They can work with measurements from the real world, or a simulator.
- ADP tend to be more data efficient - requires less data from the real world
- TD does not need to compute expectation, does not need to solve the system of linear equations - tend to be computationally more efficient.

TD vs. MC

- TD can learn/updates after every step. MC learns/updates after each episode.
- TD target depends only on one measured reward. MC target $G(s)$ depends on the sum of many rewards.
 - TD target has low variance but is biased
 - MC target is unbiased but has higher variance
- TD usually converges faster than MC in practice because TD exploits constraints of the Bellman equation

Model-based vs. Model-free

▪ Model-based RL

- First act in MDP and learn T, R
- Then value iteration or policy iteration with learned T, R
- Advantage: efficient use of data
- Disadvantage: requires building a model for T, R

▪ Model-free RL

- Bypass the need to learn T, R
- Methods to evaluate V^π , the value function for a fixed policy π without knowing T, R :
 - (i) Direct Evaluation
 - (ii) Temporal Difference Learning
- Method to learn π^*, Q^*, V^* without knowing T, R
 - (iii) Q-Learning

3. Active reinforcement learning

Actions in reinforcement learning not only gain rewards but also help learn a better model. By improving the model, greater reward may potentially be obtained. Therefore there is a need between:

- Exploitation: maximize value as reflected by the current estimate/current utility function
- Exploration: learn more about the model/model of the world to potentially improve long term well being

A scheme for balancing exploration and exploitation must:

- Try each action in each state an unbounded number of times to avoid a finite probability of missing an optimal action
- Eventually become greedy in the limit of infinite exploration

Such schemes are greedy in the limit exploration (**GILE**), e.g. ϵ -greedy is GILE if ϵ reduces to 0 at $\epsilon_k = \frac{1}{k}$.

Remark (ϵ -greedy exploration). *Choose the greedy action with probability $(1 - \epsilon)$, and an action from the policy with probability ϵ .*

Remark (ϵ). *The decreasing of ϵ is important to the convergence of an optimal value.*

3.1 Greedy action selection with $U^+(s)$

GLIE based ϵ -greedy convergence may be slow, an alternative way to balance exploration and exploitation is to use greedy action selection with respect to an **optimistic** estimate of the utility $U^+(s)$.

$$U^+(s) = R(s) + \gamma \arg \max_a f(\sum_{s'} P(s'|s, a) U^+(s'), N(s, a))$$

$N(s, a)$ is the number of times action a has been tried in state s .

3.2 Learning Q-value and Q-function

Instead of learning the utility function, we can learn an action-utility function $Q(s, a)$, the value of doing action a in state s .

Q-values are related to utility values by: $U(s) = \arg \max_a Q(s, a)$

The Q-function similarly satisfies a version of the Bellman equations:
 $Q(s, a) = R(s) + \gamma \sum_{s'} P(s'|s, a) \arg \max_{a'} Q(s', a')$

- Model-based: take ADP approach, learn $P(s' | s, a)$, then use an iterative method to compute the Q-function, given the estimated model.
 - ADP with Q-learning still need model for learning but not to take actions
 - Using MC/TD do not need model for learning and taking actions
- Model-free: an agent that learns a Q-function does not need a model of the form $P(s' | s, a)$ for action selection.

3.3 GLIE ϵ -greedy MC control

3.3.1 ϵ -greedy exploration

- Simplest idea for ensuring continual exploration
- All m actions are tried with non-zero probability / every action has a non-zero probability to be tried
- With probability ϵ choose an action at random:

$$\pi(a|s) = \begin{cases} \epsilon/m + 1 - \epsilon, & \text{if } a^* = \arg \max_{a \in A} Q(s, a) \\ \epsilon/m, & \text{otherwise} \end{cases} \quad (2)$$

ϵ -Greedy Policy Improvement:


Theorem: For any ϵ -greedy policy π , the ϵ -greedy policy π' with respect to q_π is an improvement.

3.3.2 GLIE ϵ -greedy MC control

General Idea:

- Sample k th episode using π : $\{S_1, A_1, R_1, \dots, S_T\} \sim \pi$
- For each state S_t and action A_t in the episode,
 - $N(S_t, A_t) = N(S_t, A_t) + 1$
 - $Q(S_t, A_t) = Q(S_t, A_t) + \frac{1}{N(S_t, A_t)}(G_t - Q(S_t, A_t))$
- Improve policy based on new action-value function
 - $\epsilon = \frac{1}{k}$
 - $\pi = \epsilon - \text{greedy}(Q)$

Monte-Carlo Control



Every episode:

Policy evaluation Monte-Carlo policy evaluation, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

Theorem: GLIE Monte-Carlo control converges to the optimal action-value function, $Q(s, a) \Rightarrow q_*(s, a)$

3.4 Q-learning

- With TD update, we have what is called Q-learning:

$$Q(s, a) = Q(s, a) + \alpha(R(s) + \gamma \arg \max_{a'} Q(s', a') - Q(s, a))$$
- Q-learning is **off-policy**. Works regardless of policy for generating the trajectory, e.g. random policy

3.4 SARSA

- $Q(s, a) = Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$ where a' is the action actually taken, whereas in Q-learning, a is the estimated action across different state-action pairs.
- In comparison, Q-learning uses the max over possible actions a' .
- SARSA is **on policy**. If agent policy is always exploring, learns to take exploration into account as well.
- When greedy agent that takes action with best Q-value is used, Q-learning is the same with SARSA.
- n-step SARSA

n-Step Sarsa

- Consider the following n -step returns for $n = 1, 2, \infty$:

$$\begin{array}{ll}
 n = 1 & \text{(Sarsa)} \quad q_t^{(1)} = R_{t+1} + \gamma Q(S_{t+1}) \\
 n = 2 & \quad q_t^{(2)} = R_{t+1} + \gamma R_{t+2} + \gamma^2 Q(S_{t+2}) \\
 \vdots & \quad \vdots \\
 n = \infty & \text{(MC)} \quad q_t^{(\infty)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T
 \end{array}$$

- Define the n -step Q-return

$$q_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n Q(S_{t+n})$$

- n -step Sarsa updates $Q(s, a)$ towards the n -step Q-return

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (q_t^{(n)} - Q(S_t, A_t))$$

3.5 Summary

- On-policy: given a policy, while agent follow the given policy, agent evaluate the policy and update the policy on the go.
- Off-policy: given a behavioural policy, agent evaluates other policies/target policy $\pi(a|s)$ to compute $Q_\pi(s, a)$
 - Learn from observing humans or other agents
 - Re-use experience generated from old policies π_1, π_2, \dots
 - Learn about optimal policy while following exploratory policy

- Learn about multiple policies while following one policy

- MDP & RL

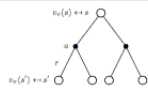
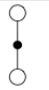
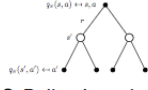
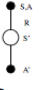
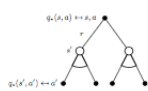

Things we know how to do:

- **If we know the MDP**
 - Compute V^* , Q^* , π^* exactly
 - Evaluate a fixed policy π
- **If we don't know the MDP**
 - We can estimate the MDP then solve
 - We can estimate V for a fixed policy π
 - We can estimate $Q^*(s,a)$ for the optimal policy while executing an exploration policy

Techniques:

- **Model-based DPs**
 - Value Iteration
 - Policy evaluation
- **Model-based RL**
- **Model-free RL**
 - Value learning
 - Q-learning

- DP & TD

	Full Backup (DP)	Sample Backup (TD)
Bellman Expectation Equation for $v_\pi(s)$	 <p>Iterative Policy Evaluation</p>	 <p>TD Learning</p>
Bellman Expectation Equation for $q_\pi(s, a)$	 <p>Q-Policy Iteration</p>	 <p>Sarsa</p>
Bellman Optimality Equation for $q_*(s, a)$	 <p>Q-Value Iteration</p>	 <p>Q-Learning</p>

4. Function approximation

Function approximation helps to scale reinforcement learning by discarding the Q-table, and generalizing from what the agent has seen to unseen. Here we focus on the **linear combinations of features** as it is differentiable

therefore easier to adjust its parameters by looking at its **gradient**.

For example, $\bar{U}_\theta(s) = \theta_1 f_1(s) + \theta_2 f_2(s) + \dots + \theta_n f_n(s)$, where $\bar{U}_\theta(s)$ is the function approximation that uses n parameters in θ to weight n features of a state to represent a very large state space. Reinforcement learning agent is supposed to learn $\theta_1, \theta_2, \dots$ to approximate the $U(s)$.

4.1. Function approximation with Monte Carlo learning

- Obtain a set of training samples: $((x_1, y_1), u_1), ((x_2, y_2), u_2), \dots, ((x_n, y_n), u_n))$, where u_j is the measured utility of the j -th example. Note, u_j is unbiased
- This gives a **supervised learning** problem.
- With squared error and linear function, we get a standard linear regression problem.
- The squared error can be minimized with **online learning**/stochastic gradient descent.
- For j -th example, the error can be represented as:

$$E_j(s) = (\bar{U}_\theta(s) - u_j(s))^2 \quad (3)$$

- i -th parameter θ_i can be updated as:

$$\begin{aligned} \theta_i(s) &= \theta_i \underbrace{-}_{\text{Negative gradient}} \alpha \frac{\partial E_j}{\partial \theta_i} \\ &= \theta_i + \alpha (\bar{U}_\theta(s) - u_j(s)) \frac{\partial \bar{U}_\theta(s)}{\partial \theta_i} \end{aligned} \quad (4)$$

4.2 Online learning with temporal difference learning :

- Obtain a set of training data: $((s_1, (R(s_2) + \gamma U_\theta(s_2))), (s_2, (R(s_3) + \gamma U_\theta(s_3))), \dots, (s_n, (R(s_{n+1}) + \gamma U_\theta(s_{n+1}))))$

- For TD: $\theta_i = \theta_i + \alpha(R(s) + \gamma \bar{U}_\theta(s') - \bar{U}_\theta(s)) \frac{\partial \bar{U}_\theta(s)}{\partial \theta_i}$
- For Q-learning: $\theta_i = \theta_i + \alpha(R(s) + \gamma \arg \max_{a'} \bar{Q}_\theta(s', a') - \bar{Q}_\theta(s, a)) \frac{\partial \bar{Q}_\theta(s, a)}{\partial \theta_i}$

5. Policy search

5.1. Policy representation and search

Policy search adjusts θ to improve the policy. Policy search tries to find a policy, e.g. represented as Q-functions that does well, so $Q^*/10$ can give the same optimal actions as Q^* . In contrast, Q-learning with function approximation tries to find a value of θ such that \bar{Q}_θ that is close to Q^* .

- Use **stochastic policy** $\pi_\theta(s, a)$ that specifies the probability of selecting action a in state s . This solves the problem that policy as a function of action is discontinuous. For example, the **softmax function**:

$$\underbrace{\pi_\theta(s, a)}_{\text{Probability distribution of actions in } s} = \frac{e^{\bar{Q}_\theta(s, a)}}{\underbrace{\sum_{a'} e^{\bar{Q}_\theta(s, a')}}_{\text{Nomalization}}}$$

- We can specify the **policy value** as $\rho(\theta)$. $\rho(\theta)$ can be optimized by:
 - Taking a step in the direction of the **policy gradient** $\nabla_\theta \rho(\theta)$ /hill climbing (positive gradient descent), if $\rho(\theta)$ is differentiable (if we specify the policy in softmax function).
 - Look for a local optimal
- For stochastic environment and/or policy $\pi_\theta(s, a)$, it is possible to obtain an unbiased estimate of gradient at θ , $\nabla_\theta \rho(\theta)$ directly from results of trials executed at θ .
- Consider **single action from single state** s_0 :

$$\nabla_\theta \rho(\theta) = \nabla_\theta \sum_a \pi_\theta(s_0, a) R(a) = \sum_a \nabla_\theta (\pi_\theta(s_0, a)) R(a)$$

Then we approximate the summation using **samples** generated from $\pi_\theta(s_0, a)$ / sample N samples (a_n) from one state:

$$\nabla_\theta \rho(\theta) = \sum_a \pi_\theta(s_0, a) \frac{\nabla_\theta(\pi_\theta(s_0, a))R(a)}{\pi_\theta(s_0, a)} \approx \frac{1}{N} \sum_{j=1}^n \frac{\nabla_\theta(\pi_\theta(s_0, a_j))R(a_j)}{\pi_\theta(s_0, a_j)}$$

5.2. Policy search: REINFORCE algorithm

- From **single state**, we generalize this idea to **sequential case**: state s_i and action a_{ij} , $\pi_\theta(s_i, a_{ij})$.

$$\nabla_\theta \rho(\theta) \propto \sum_s \rho_{\pi_\theta}(s) \sum_a \frac{\pi_\theta(s_0, a) \nabla_\theta \pi_\theta(s_0, a) Q_{\pi_\theta}(s, a)}{\pi_\theta(s_0, a)} \approx \frac{1}{N} \sum_i^N \sum_j \frac{\nabla_\theta \pi_\theta(s_i, a_{ij}) G_{ij}(s_i)}{\pi_\theta(s_i, a_{ij})},$$

for each state s_i is visited, a_{ij} is executed on j -th trial, and G_{ij} is the total reward/return received from state s_i onwards on j -th trial.

- Use online learning/update, we get **REINFORCE** algorithm:

$$\theta_{j+1} = \theta_j + \alpha G_j \frac{\nabla_\theta \pi_\theta(s, a_j)}{\pi_\theta(s, a_j)}$$

As $\nabla_\theta \ln(\pi_\theta(s, a_j)) = \frac{\nabla_\theta \pi_\theta(s, a_j)}{\pi_\theta(s, a_j)}$, we have rewrite our update function as follows:

$\theta_{j+1} = \theta_j + \alpha G_j \nabla_\theta \ln(\pi_\theta(s, a_j))$, because usually the gradient of natural logarithm of the policy function is easier to look for than directly from the policy function itself.

function REINFORCE

Initialise θ arbitrarily

for each episode $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$ **do**

for $t = 1$ to $T - 1$ **do**

$\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$

end for

end for

return θ

end function

- Reduce variance with a **Baseline**

We are estimating $\nabla \rho(\theta) = \sum_s p_{\pi_\theta}(s) \sum_a \nabla_\theta \pi_\theta(s, a_j) Q_{\pi_\theta}(s, a)$

$$\begin{aligned}
\nabla \rho(\theta) &= \sum_s p_{\pi_\theta}(s) \sum_a \nabla_\theta \pi_\theta(s, a_j) Q_{\pi_\theta}(s, a) \\
&= \sum_s p_{\pi_\theta}(s) \sum_a \nabla_\theta \pi_\theta(s, a_j) (Q_{\pi_\theta}(s, a) - B(s)) \\
&= \sum_s p_{\pi_\theta}(s) \left(\sum_a \nabla_\theta \pi_\theta(s, a_j) Q_{\pi_\theta}(s, a) - \left(\sum_a \nabla_\theta \pi_\theta(s, a_j) B(s) \right) \right) \\
&= \sum_s p_{\pi_\theta}(s) \left(\sum_a \nabla_\theta \pi_\theta(s, a_j) Q_{\pi_\theta}(s, a) - B(s) \underbrace{\left(\sum_a \nabla_\theta \pi_\theta(s, a_j) \right)}_{\text{Distribution sums up to 1}} \right) \\
&= \sum_s p_{\pi_\theta}(s) \left(\sum_a \nabla_\theta \pi_\theta(s, a_j) Q_{\pi_\theta}(s, a) - B(s) (\nabla_\theta 1) \right)
\end{aligned} \tag{5}$$

- Using a baseline function $B(s)$ can reduce variance
- Use an **advantage function** in place of $Q_{\pi_\theta}(s, a)$: $A_{\pi_\theta}(s, a) = Q_{\pi_\theta}(s, a) - V_{\pi_\theta}(s)$, where $V_{\pi_\theta}(s)$ is the baseline function
- REINFORCE uses a Monte Carlo estimate of the advantage function, which has higher variance. To reduce variance, an alternative is to use TD method. The advantage function is: $Q_{\pi_\theta}(s, a) - V_{\pi_\theta}(s) = E[r + \gamma V_{\pi_\theta}(s')] - V_{\pi_\theta}(s)$. It is also common to use multiple steps of rewards instead of one step in TD.
- **actor-critic method** :
 - **Critic**: Learns a value or Q-function, that is to update parameter w , that is used only for evaluation
 - **Actor**: Learns a policy (actor) that takes action, that is to update parameters θ in direction suggested by critic
 - The critic is solving a familiar problem: policy evaluation

- We use a critic to estimate the action-value function, $Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$

Algorithm 1: Adaptive Dynamic Programming

```

1 function PASSIVE-ADP-AGENT(percept) returns an action ;
   Input   : percept, a percept indicating the current state  $s'$  and reward
              signal  $r'$ 
   Output:  $a$ 
2 Persistent:  $\pi$ , a fixed policy; mdp, an MDP with model P, reward R,
              discount  $\gamma$ ; U, a table of utilities, initially empty;  $N_{sa}$ , a table of
              frequencies for state-action pairs, initially zero;  $N_{s'|sa}$ , a table of
              outcome frequencies given state-action pairs, initially zero; s, a, the
              previous state and action, initially null ;
3 if  $s'$  is new then
4   |  $U[s'] = r'$  ;
5   |  $R[s'] = r'$ ;
6 end
7 if  $s$  is not null then
8   | increment  $N_{sa}[s, a]$  and  $N_{s'|sa}[s', s, a]$  ;
9   | foreach  $t$  such that  $N_{t|sa}[t, s, a]$  is nonzero do
10    |  $P(t|s, a) = N_{s'|sa}[s', s, a] / N_{sa}[s, a]$ ;
11    | end
12 end
13 U = POLICY-EVALUATION ( $\pi$ , U, mdp);
14 if  $s'$  is Terminal then
15   | s, a = null ;
16   |  $R[s'] = r'$ ;
17 else
18   | s, a =  $s', \pi[s']$ 
19 end
20 return a

```

Algorithm 2: Temporal-Difference Learning

```
1 function PASSIVE-TD-AGENT(percept) returns an action ;  
   Input : percept, a percept indicating the current state  $s'$  and reward  
          signal  $r'$   
   Output:  $a$   
2 Persistent:  $\pi$ , a fixed policy; mdp, an MDP with model P, reward R,  
   discount  $\gamma$ ; U, a table of utilities, initially empty;  $N_{sa}$ , a table of  
   frequencies for state-action pairs, initially zero; s, a, r, the previous  
   state, action, and reward, initially null ;  
3 if  $s'$  is new then  
4   |  $U[s'] = r'$  ;  
5 end  
6 if  $s$  is not null then  
7   | increment  $N_{sa}[s, a]$  ;  
8   |  $U^\pi[s] = U^\pi[s] + \alpha(N_{sa}[s])(R(s) + U\pi[s'] - U\pi[s])$   
9 end  
10 if  $s'$  is Terminal then  
11   |  $s, a, r = \text{null}$  ;  
12 else  
13   |  $s, a, r = s', \pi[s'], r$   
14 end  
15 return a
```

Algorithm 3: Q Learning

```
1 function Q-LEARNING-AGENT(percept) returns an action ;  
   Input : percept, a percept indicating the current state  $s'$  and reward  
           signal  $r'$   
   Output:  $a$   
2 Persistent:  $Q$ , a table of action values indexed by state and action,  
   initially zero;  $N_{sa}$ , a table of frequencies for state-action pairs,  
   initially zero;  $s$ ,  $a$ ,  $r$ , the previous state, action, and reward, initially  
   null ;  
3 if  $s'$  is Terminal then  
4   |  $Q[s', \text{None}] = r'$   
5 end  
6 if  $s$  is not null then  
7   | increment  $N_{sa}[s, a]$  ;  
8   |  $Q(s, a) = Q(s, a) + \alpha(N_{sa}[s, a])(R(s) + \gamma \arg \max_{a'} Q(s', a') - Q(s, a))$   
9 end  
10  $s, a, r = s', \arg \max_{a'} f(Q(s', a'), N(s', a')), r'$  ;  
11 return  $a$ 
```
